

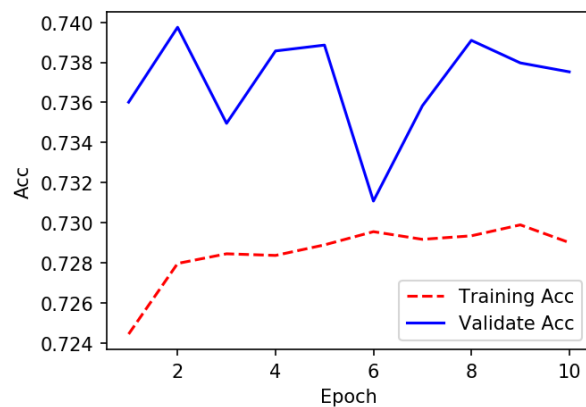

Model summary is followed:

Layer (type)	Output Shape	Param #
word_embedding_layer (Embedd	(None, 23, 300)	27914700
flatten_2 (Flatten)	(None, 6900)	0
hidden_layer (Dense)	(None, 128)	883328
output_layer (Dense)	(None, 2)	258

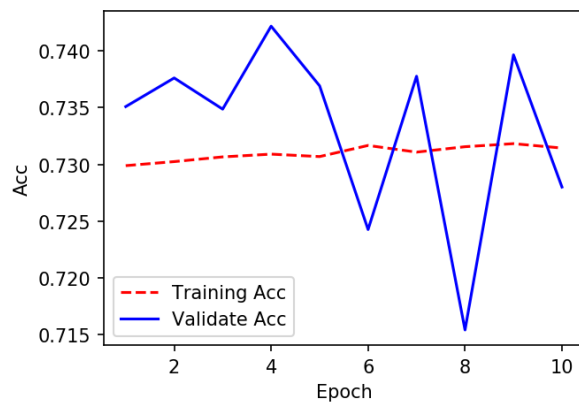
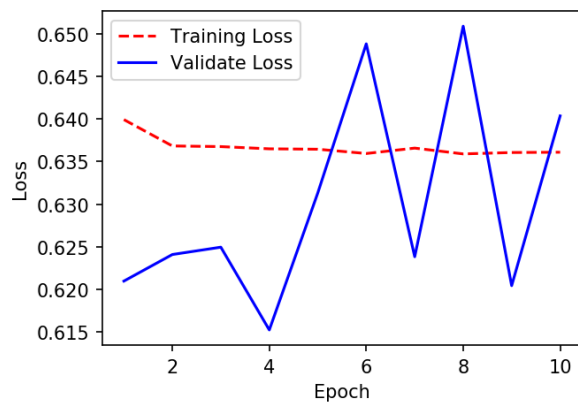
Total params: 28,798,286
Trainable params: 883,586
Non-trainable params: 27,914,700

Dropout rate is 0.3 and. L2-norm regularization is applied. Trained models as below.



relu activation dropout_0.3 and L2-norm regularization



Tanh Activation dropout_0.3 and L2



Sigmoid Activation dropout_0.3 and L2-norm Regularization



Activation Function	Accuracy (%)
Tanh	76.33
Sigmoid	77.64
Relu	77.65


ReLU is less computationally expensive than tanh and sigmoid because it involves simpler mathematical operations. That is a good point to consider when we are designing deep neural nets. From the results, the activation functions have not influenced the accuracy significantly.

What effect does addition of L2-norm regularization have on the results?

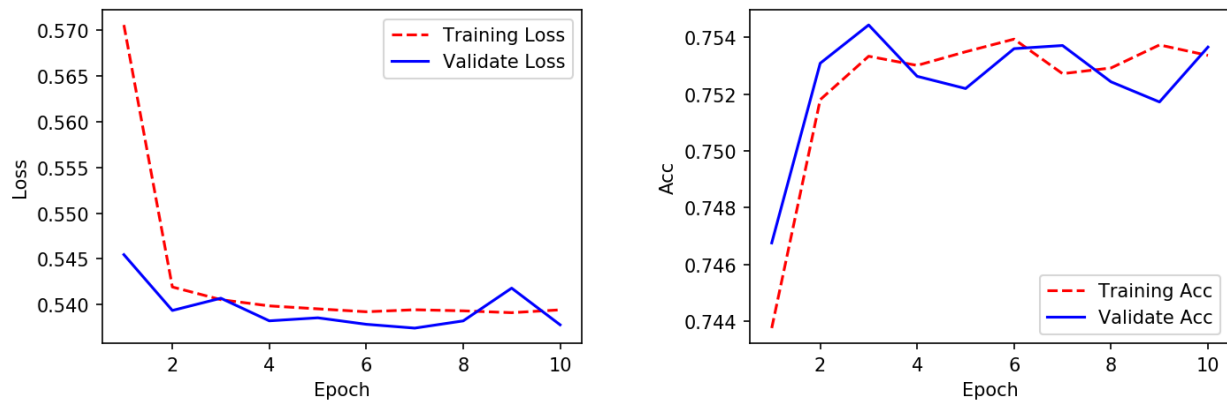
Without L2-norm regularization, the accuracy of the validation dataset is going up, and then dropping down, which means that the model is overfitted. As the figures followed show that the L2 regularized model has become much more resistant to overfitting than the baseline model, even

though both models have the same number of parameters. $12 (0.001)$ means that every coefficient in the weight matrix of the layer will add $0.001 * \text{weight_coefficient_value}^{**2}$ to the total loss of the network.

relu activation



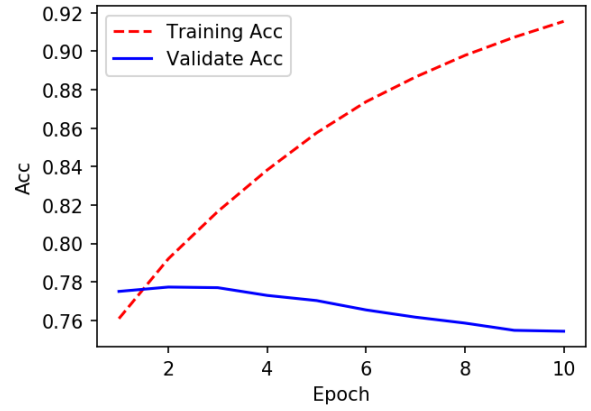
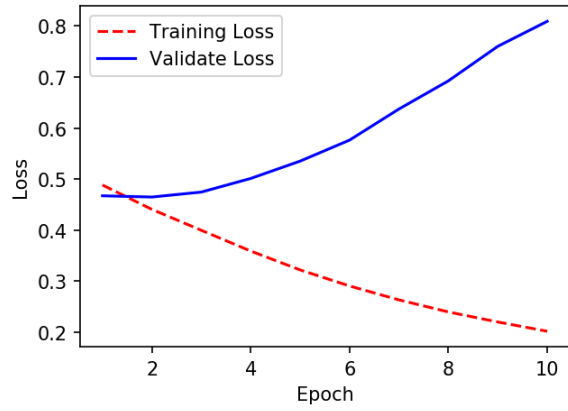
Relu Activation L2 norm Regularization





What effect does dropout have on the results?

Dropout is also one of the most effective and most commonly used regularization techniques for neural networks. After applying dropout, the overfitting is prevented by the appropriate dropout rate, which is the fraction of the features that are being zeroed-out; it is usually set between 0.2 and 0.5. From the figures below, when the dropout rate is 0.1, the overfitting cannot be prevented, while the dropout rate of 0.3 and 0.6 work well.

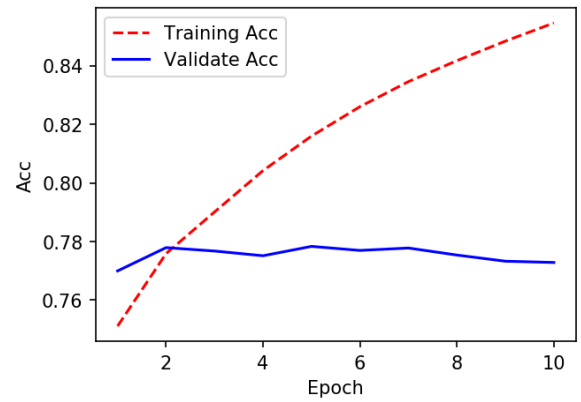
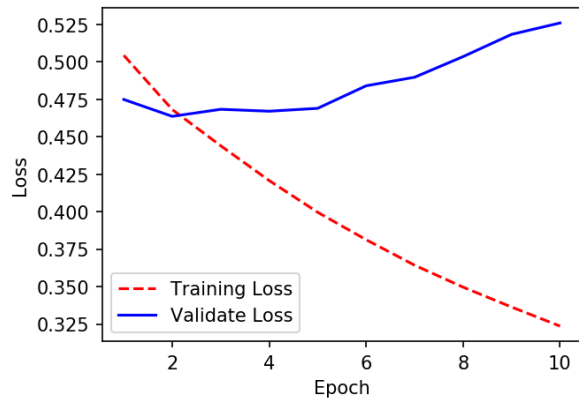
relu activation



Relu Activation dropout_0.1



relu activation dropout_0.3



Relu Activition dropout_0.6

